

Approximated PCA

Rodrigo Arias

February 17, 2017

PCA algorithm

1. Take a dataset X of n variables.
2. Scale and center the variables.
3. From X compute the $n \times n$ covariance matrix S .
4. **Compute the eigenvalues and eigenvectors of S .**
5. *Optional: Ignore some eigenvectors.*
6. Generate a new basis from the selected eigenvectors.
7. Project X into the new basis.

PCA algorithm

1. Take a dataset X of n variables.
2. Scale and center the variables.
3. From X compute the $n \times n$ covariance matrix S .
4. **Compute the eigenvalues and eigenvectors of S .**
5. *Optional: Ignore some eigenvectors.*
6. Generate a new basis from the selected eigenvectors.
7. Project X into the new basis.

Computing the eigenvectors is the principal step of PCA.

Visual example

The input dataset with $n = 800$ variables and 1023 observations (columns),

$$X = \begin{pmatrix} x_{1,1} & x_{1,2} & \dots & x_{1,1023} \\ x_{2,1} & x_{2,2} & \dots & x_{2,1023} \\ \vdots & \vdots & & \vdots \\ x_{800,1} & x_{800,2} & \dots & x_{800,1023} \end{pmatrix}$$

Each pixel $x_{i,j}$ is a value in $[0, 255]$ of the gray level in the image.

Visual example

The input dataset with $n = 800$ variables and 1023 observations (columns), $X =$



Each pixel $x_{i,j}$ is a value in $[0, 255]$ of the gray level in the image.


Visual example

Keep 800 eigenvectors. Output:



Visual example

Keep 200 eigenvectors. Output:



Visual example

Keep 100 eigenvectors. Output:



Visual example

Keep 40 eigenvectors. Output:



Computing eigenvalues and eigenvectors

1. Take the $n \times n$ target matrix $A = S$.
2. Compute a **tridiagonal** matrix T : $A = PTP^T$.
3. From T compute a **diagonal** matrix D : $T = QDQ^T$.
4. The eigenvalues of A are in the diagonal of D .
5. Compute the eigenvectors from D .

Computing eigenvalues and eigenvectors

1. Take the $n \times n$ target matrix $A = S$.
2. Compute a **tridiagonal** matrix T : $A = PTP^T$.
3. From T compute a **diagonal** matrix D : $T = QDQ^T$.
4. The eigenvalues of A are in the diagonal of D .
5. Compute the eigenvectors from D .

Computing a tridiagonal matrix is called **tridiagonalization**. For the diagonal, **diagonalization**. Several algorithms exists for both steps.

Tridiagonalization

These algorithms transform a **symmetric** matrix A into a new pair of matrices P and T such that P is orthogonal, T is tridiagonal, and $A = PTP^T$

$$\begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{pmatrix} = P \begin{pmatrix} t_{11} & t_{12} & & \\ t_{21} & t_{22} & t_{23} & \\ & t_{32} & t_{33} & t_{34} \\ & & t_{43} & t_{44} \end{pmatrix} P^T$$

Tridiagonalization

Algorithm	Complexity	Iterative	Stability
Householder	$O(4n^3/3)$	No	Great
Givens	$O(kn^3)$	No	Good
Lanczos	$O(kpn^2)$	Yes	Bad
Others			

Where $n \times n$ is the dimension of the matrix A , k is some constant, and p the number of iterations.

Diagonalization

These algorithms take a **tridiagonal** matrix T into a new pair of matrices Q and D such that Q is orthogonal, D is diagonal, and $T = QDQ^T$

$$\begin{pmatrix} t_{11} & t_{12} & & \\ t_{21} & t_{22} & t_{23} & \\ & t_{32} & t_{33} & t_{34} \\ & & t_{43} & t_{44} \end{pmatrix} = Q \begin{pmatrix} d_{11} & & & \\ & d_{22} & & \\ & & d_{33} & \\ & & & d_{44} \end{pmatrix} Q^T$$

The matrix D contains the **eigenvalues** in the diagonal.

Diagonalization

Algorithm	Complexity	Convergence
QR	$O(6n^3)$	Cubic
Divide and conquer	$O(8n^3/3)$	Quadratic
Jacobi	$O(n^3)$	Quadratic
Power iteration	$O(n^3)$	Linear
Inverse iteration	$O(n^3)$	Linear
Others		

All these algorithms are iterative and $n \times n$ is the dimension of the matrix A .

The selected method

Householder + QR

Approximated computing

Maximize an utility function, based on your preferences:

$$\max u(E, T, n, \epsilon, C, \dots)$$

E Energy consumption

T Time taken

n Problem size

ϵ Error in the result

C Economical cost

$u()$ Utility function

Techniques: Deterministic

- ▶ **Bit-Width Reduction**
- ▶ Float-to-Fixed Conversion
- ▶ Fuzzy Memoization
- ▶ Hierarchical FPU
- ▶ Load Value Approximation
- ▶ Lossy Compression and Data
- ▶ Packing
- ▶ Precision Scaling (ALU)
- ▶ Reduced-Precision FPU
- ▶ Underdesigned Multiplier
- ▶ ...

Techniques: Deterministic

- ▶ **Bit-Width Reduction**
- ▶ Float-to-Fixed Conversion
- ▶ Fuzzy Memoization
- ▶ Hierarchical FPU
- ▶ Load Value Approximation
- ▶ Lossy Compression and Data
- ▶ Packing
- ▶ Precision Scaling (ALU)
- ▶ Reduced-Precision FPU
- ▶ Underdesigned Multiplier
- ▶ ...

Begin with bit-width reduction

Bit-width reduction in Householder

- ▶ The Householder algorithm performs $n - 2$ steps.
- ▶ In each step, a matrix multiplication is done.
- ▶ With less precision, the error in the final result can grow.
- ▶ Analyze how the error changes in Householder.

Bit-width reduction in Householder

- ▶ Some theoretical bounds can be used to check the experimental error
- ▶ Design an emulator with an ALU of b bit-width
- ▶ Test the algorithm in the emulator with different b , and check the error.